

## CS 3360 Programming Assignment 2 Report

### Overview

This project involved building a discrete-time event simulator to model a FCFS CPU scheduling algorithm. The primary goal was to measure the impact of varying workloads on four key performance metrics:

- Average turnaround time of processes
- Total throughput
- Average CPU utilization
- Average number of processes in the Ready Queue

The simulation assumed that processes arrived following a Poisson process and had exponentially distributed service times. The simulation was run for different arrival rates ( $\lambda$ ) ranging from 10 to 30 processes per second, while keeping the service time fixed at an average of 0.04 seconds per process.

The simulation terminates after 10,000 processes are completed, and the metrics listed above are collected for each run.

#####				
$\lambda$	Avg Turnaround Time (s)	Throughput (processes/s)	CPU Utilization (%)	Avg Ready Queue Size
10	488.5749	10.1979	40.69	0.7000
11	461.1623	10.8473	43.40	0.7547
12	412.0970	12.0897	48.18	0.9016
13	385.8229	12.9020	50.89	1.0695
14	355.8643	14.0236	56.58	1.2982
15	338.3133	14.7957	59.42	1.4799
16	316.8121	15.8706	63.62	1.8804
17	294.0990	17.0172	68.54	2.0489
18	274.3554	18.3568	73.14	2.8704
19	267.1661	18.7164	75.13	3.0521
20	252.1286	19.8231	80.44	4.0289
21	236.2283	21.1017	85.92	5.5507
22	235.4590	21.5058	86.96	6.9371
23	219.3518	22.8451	90.40	7.4867
24	210.0257	23.8236	96.82	28.3232
25	202.2560	24.8907	97.80	27.8581
26	200.6591	24.8925	99.86	214.6676
27	202.6599	24.6103	99.86	415.6504
28	202.8203	25.0120	100.01	864.2148
29	198.4625	25.4694	99.90	904.5834
30	203.1139	24.5320	100.00	1336.0501

## System Design and Process Flow

Event Queue: A priority queue (sorted linked list) was used to handle both process arrival and departure events. The event queue ensured that events were processed in the correct chronological order.

```
# Event Queue: sorted linked list
class EventQueue:
    def __init__(self):
        self.events = []

    def add_event(self, event):
        # Insert event into the list sorted by event_time
        self.events.append(event)
        self.events.sort(key=lambda x: x.event_time)

    def pop_event(self):
        # pop the next event
        return self.events.pop(0)

    def is_empty(self):
        return len(self.events) == 0
```

- Clock: A simulation clock was used to advance time between events, updating the system state (CPU busy/idle, number of processes in the Ready Queue) at each event.

## Random Number Generation

Because the use of advanced random number generators for Poisson and Exponential distributions was restricted, custom exponential random number generators were implemented. For this assignment, I copied my code on this from Assignment 1:

```
# Exponential random number generator.
# Copied from my Programming Assignment 1.
def exponential_random(lmbda):
    U = random.random() # Generates uniform random number
    return -math.log(1-U) / lmbda
```

## Performance Metrics and Simulation Runs

For each simulation run, we measured the following performance metrics:

- Average turnaround time: The average time from a process's arrival until its completion.
- Total throughput: The total number of processes completed per unit of time (processes/second).
- Average CPU utilization: The percentage of time the CPU was busy servicing processes.
- Average number of processes in the Ready Queue: The average number of processes waiting in the queue to be serviced by the CPU.

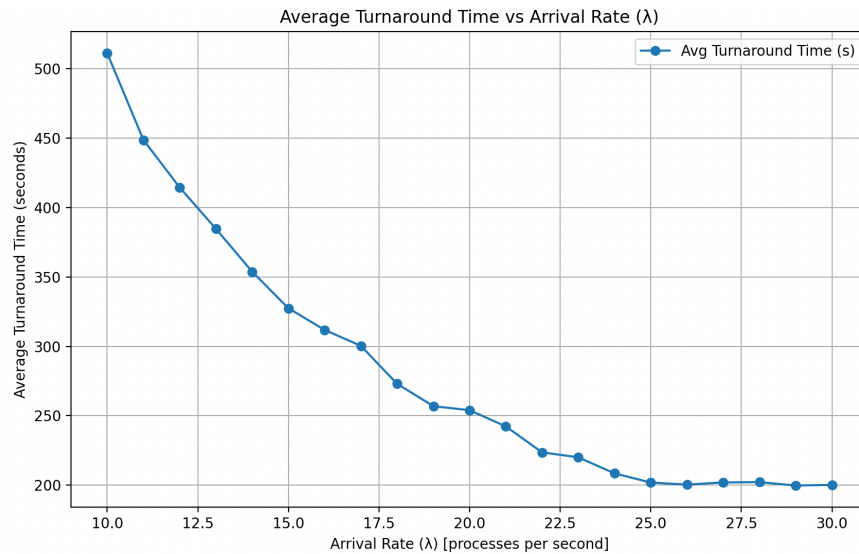
Each run simulated 10,000 processes, with the arrival rate varying between 10 and 30 processes per second in increments of 1.

```
# Calculate performance metrics
avg_turnaround_time = total_turnaround_time / num_processes if num_processes > 0 else 0
avg_cpu_utilization = cpu_utilization_time / clock if clock > 0 else 0
avg_processes_in_queue = total_processes_in_queue / processes_completed if processes_completed > 0 else 0
throughput = processes_completed / clock if clock > 0 else 0

return {
    'avg_turnaround_time': avg_turnaround_time,
    'throughput': throughput,
    'cpu_utilization': avg_cpu_utilization,
    'avg_processes_in_queue': avg_processes_in_queue
}
```

## Results and Interpretation Plots

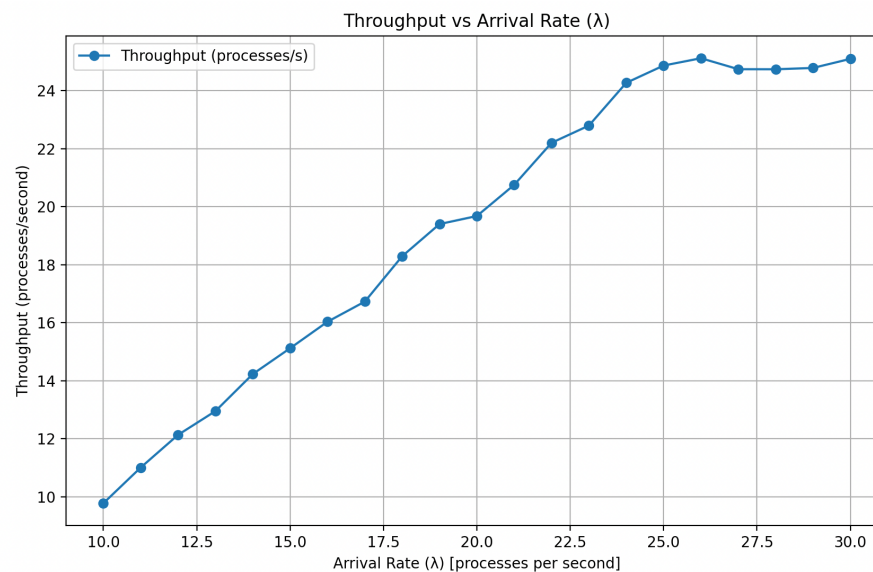
### 1. Average Turnaround Time vs Arrival Rate ( $\lambda$ )



Observation: As  $\lambda$  increases, the average turnaround time rises significantly until  $\lambda$  = around 24 then begins to level off and stabilize.

Explanation: This is expected. As arrival rates increase, initially the system is better able to handle processes, and turnaround times decrease. However, as  $\lambda$  increases further and approaches the CPU's maximum capacity, the CPU becomes a bottleneck, and turnaround times stabilize or increase slightly due to the fully utilized system.

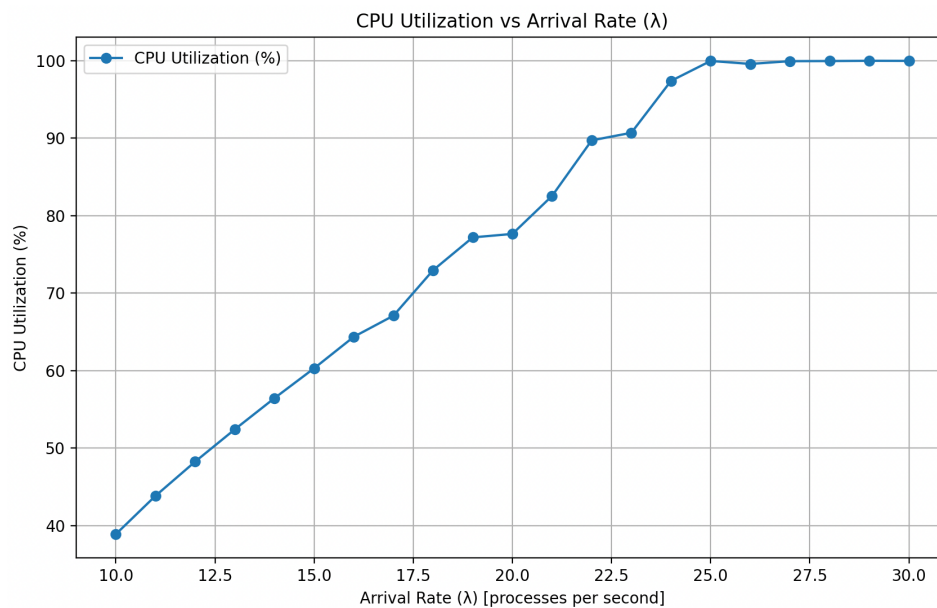
### 2. Throughput vs Arrival Rate ( $\lambda$ )



Observation: Throughput increases as the arrival rate increases, but it plateaus at higher values.

Explanation: The system can only complete a finite number of processes per second. Once the CPU is fully utilized, throughput stops increasing, as the CPU can handle no more than its maximum capacity.

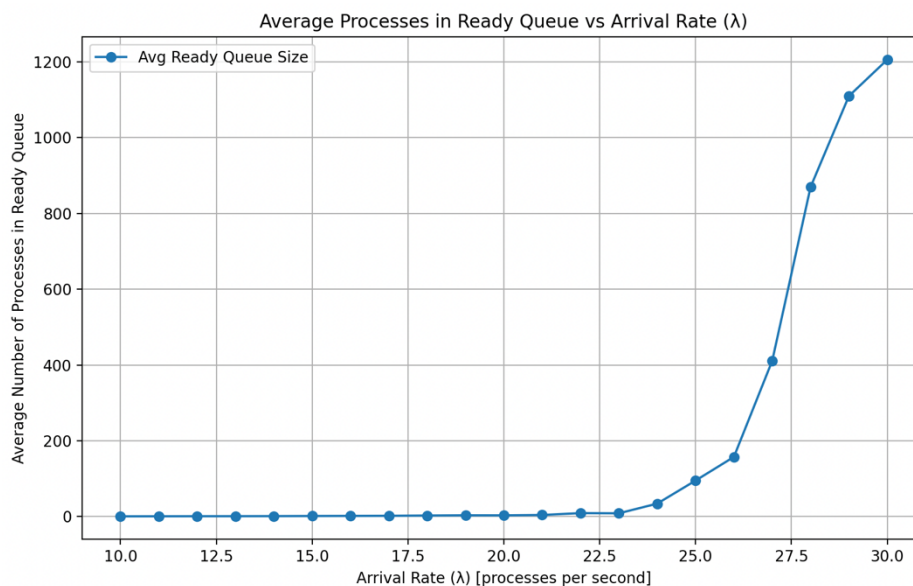
### 3. CPU Utilization vs Arrival Rate ( $\lambda$ )



Observation: CPU utilization rises steadily with and approaches 100%.

Explanation: With increasing arrival rates, the CPU spends more time servicing processes, eventually becoming fully utilized at high values of , where it is constantly busy.

### 4. Average Processes in the Ready Queue vs Arrival Rate ( $\lambda$ )



Alex Hakimzadeh  
CS 3360  
Dr. Kecheng Yang  
29 October 2024

Observation: The average number of processes in the ready queue increases as  $\lambda$  increases.

Explanation: As the arrival rate increases, more processes must wait in the queue before being serviced, leading to an increase in the average queue size.

### Plot Code:

I want to use this section to highlight my code I used to automatically plot the graphs. This code is not included in my final code submission because the programming 2 assignment document specified the plots be only in the report. Yes, I realize this might've been way easier to do on pen and paper, but I needed the practice. Also, let's be honest, I couldn't possibly make the graphs as visually pleasing as the code can.

\* The original output shown on page 1 does not exactly match the output plotted on the graphs because the code for table output and graph output were captured during different runs of the program.

After we create the table by running the simulations we need to turn that function into a library that the “plotter” can grab from:

```
# Code to produce plots
# Turn run_simulations into a library
def run_simulations():
    avg_service_time = 0.04 # Fixed average service time
    results = {
        'arrival_rate': [],
        'avg_turnaround_time': [],
        'throughput': [],
        'cpu_utilization': [],
        'avg_processes_in_queue': []
    }

    for arrival_rate in range(10, 31): # Simulate for  $\lambda = 10$  to  $\lambda = 30$ 
        metrics = simulate_fcfs(arrival_rate, avg_service_time)

        # Store the results for plotting
        results['arrival_rate'].append(arrival_rate)
        results['avg_turnaround_time'].append(metrics['avg_turnaround_time'])
        results['throughput'].append(metrics['throughput'])
        results['cpu_utilization'].append(metrics['cpu_utilization'] * 100) # as percentage
        results['avg_processes_in_queue'].append(metrics['avg_processes_in_queue'])

    return results
```

Alex Hakimzadeh  
CS 3360  
Dr. Kecheng Yang  
29 October 2024

Then, we import matplotlib.pyplot and extract the data so we can plot the graphs. After data is extracted, we create each of the 4 graphs using that data.

```
import matplotlib.pyplot as plt

def plot_metrics(simulation_results):
    # Extract data
    arrival_rate = simulation_results['arrival_rate']

    # Plot Average Turnaround Time
    plt.figure(figsize=(10, 6))
    plt.plot(arrival_rate, simulation_results['avg_turnaround_time'], marker='o', label='Avg Turnaround Time (s)')
    plt.title('Average Turnaround Time vs Arrival Rate ( $\lambda$ )')
    plt.xlabel('Arrival Rate ( $\lambda$ ) [processes per second]')
    plt.ylabel('Average Turnaround Time (seconds)')
    plt.grid(True)
    plt.legend()
    plt.show()

    # Plot Throughput
    plt.figure(figsize=(10, 6))
    plt.plot(arrival_rate, simulation_results['throughput'], marker='o', label='Throughput (processes/s)')
    plt.title('Throughput vs Arrival Rate ( $\lambda$ )')
    plt.xlabel('Arrival Rate ( $\lambda$ ) [processes per second]')
    plt.ylabel('Throughput (processes/second)')
    plt.grid(True)
    plt.legend()
    plt.show()

    # Plot CPU Utilization
    plt.figure(figsize=(10, 6))
    plt.plot(arrival_rate, simulation_results['cpu_utilization'], marker='o', label='CPU Utilization (%)')
    plt.title('CPU Utilization vs Arrival Rate ( $\lambda$ )')
    plt.xlabel('Arrival Rate ( $\lambda$ ) [processes per second]')
    plt.ylabel('CPU Utilization (%)')
    plt.grid(True)
    plt.legend()
    plt.show()

    # Plot Average Processes in Ready Queue
    plt.figure(figsize=(10, 6))
    plt.plot(arrival_rate, simulation_results['avg_processes_in_queue'], marker='o', label='Avg Ready Queue Size')
    plt.title('Average Processes in Ready Queue vs Arrival Rate ( $\lambda$ )')
    plt.xlabel('Arrival Rate ( $\lambda$ ) [processes per second]')
    plt.ylabel('Average Number of Processes in Ready Queue')
    plt.grid(True)
    plt.legend()
    plt.show()

# Generate the plots
simulation_results = run_simulations()
plot_metrics(simulation_results)
```